



# LAMMPS on GPUs

## *A Tutorial*

W. Michael Brown, Peng Wang, Paul S. Crozier, Steve  
Plimpton

Wednesday, February 24, 2010

# Why run on GPUs?

- Technology paid for by gamers, but impact to scientific computing is now well-recognized
- Cheap, low-power (electrical) solution for data parallelism
  - 240+ cores on a GPU
  - High memory bandwidth



# Porting LAMMPS to GPUs


- Still largely a research effort

**Marc Adams (Nvidia)**  
**Pratul Agarwal (ORNL)**  
**Sarah Anderson (Cray)**  
**Mike Brown (Sandia)**  
**Paul Crozier (Sandia)**  
**Massimiliano Fatica (Nvidia)**  
**Scott Hampton (ORNL)**  
**Ricky Kendall (ORNL)**  
**Hyesoon Kim (Ga Tech)**  
**Axel Kohlmeyer (Temple)**  
**Doug Kothe (ORNL)**  
**Scott LeGrand (Nvidia)**

**Ben Levine (Temple)**  
**Christian Mueller (UTI Germany)**  
**Steve Plimpton (Sandia)**  
**Duncan Poole (Nvidia)**  
**Steve Poole (ORNL)**  
**Jason Sanchez (RPI)**  
**Arnold Tharrington (ORNL)**  
**John Turner (ORNL)**  
**Peng Wang (Nvidia)**  
**Lars Winterfeld (UTI Germany)**  
**Andrew Zonenberg (RPI)**

# Currently Available in “Main” LAMMPS

- Lennard-Jones
  - Force/Neighbor
- Gay-Berne Potential
  - Force
- More capabilities soon...



# How to Run LAMMPS on Your GPU

# 1. Do you have a GPU?

- For single precision
  - Currently need a CUDA-enabled GPU with compute capability  $\geq 1.1$
- For double precision
  - Currently need a CUDA-enabled GPU with compute capability  $\geq 1.3$

Windows: Device Manager

Apple: “Apple Menu”-> “About this Mac” -> “More Info” -> “Graphics/Displays”

Linux: nvidia\_settings or /sbin/lspci | grep nVidia

List of CUDA-enabled GPUs here:

[http://www.nvidia.com/object/cuda\\_gpus.html](http://www.nvidia.com/object/cuda_gpus.html)

Can use device query to get compute capability; more later...

## 2. Do you have CUDA?

- [http://developer.nvidia.com/object/cuda\\_2\\_3\\_downloads.html](http://developer.nvidia.com/object/cuda_2_3_downloads.html)
  - Need driver and toolkit only
  - Need to have the nvcc compiler in your path
  - Pay attention to 32- or 64-bit
    - No 64-bit on apple!

```
set path = ( $path /usr/local/cuda/bin )
setenv LD_LIBRARY_PATH /usr/local/cuda/lib/
or
set path = ( $path /usr/local/cuda/bin )
setenv LD_LIBRARY_PATH /usr/local/cuda/lib64/
```

# 3. Edit LAMMPS GPU Makefile

```
set LROOT = "/home/wmbrown/lammps-20Feb10"  
cd $LROOT/lib/gpu  
emacs Makefile.nvidia
```

# 3. Edit LAMMPS GPU Makefile (2)

```
BIN_DIR = .
OBJ_DIR = .
AR = ar
CUDA_CPP = nvcc -I/usr/local/cuda/include -DUNIX -O3 -Xptxas -v --
use_fast_math
CUDA_ARCH = -arch=sm_13
CUDA_PREC = -D_SINGLE_SINGLE
CUDA_LINK = -L/usr/local/cuda/lib64 -lcudart $(CUDA_LIB)
```

For compute capability >= 1.3 can also use:

```
CUDA_PREC = -D_SINGLE_DOUBLE # Double precision accumulation
or
CUDA_PREC = -D_DOUBLE_DOUBLE # Double precision everything
```

For Apple, must compile 32-bit

```
CUDA_ARCH = -arch=sm_13 -m32
CUDA_LINK = -L/usr/local/cuda/lib -lcudart $(CUDA_LIB)
```

For compiler >= g++ 4.4 on Linux

```
CUDA_ARCH = -arch=sm_13 --compiler-bindir=/usr/bin/gcc-4.3
```

# 4. Make LAMMPS GPU lib

```
make -f Makefile.nvidia  
./nvc_get_devices
```

Device 0: "GeForce GTX 295"	
Revision number:	1.3
Total amount of global memory:	0.87 GB
Number of multiprocessors:	30
Number of cores:	240
Total amount of constant memory:	65536 bytes
Total amount of shared memory per block:	16384 bytes
Total number of registers available per block:	16384
Warp size:	32
Maximum number of threads per block:	512
Maximum sizes of each dimension of a block:	512 x 512 x 64
Maximum sizes of each dimension of a grid:	65535 x 65535 x 1
Maximum memory pitch:	262144 bytes
Texture alignment:	256 bytes
Clock rate:	1.24 GHz
Concurrent copy and execution:	Yes
Device 1: "Tesla C1060"	

# 5. Edit LAMMPS Makefile as Necessary

```
cd $LROOT/src  
emacs ./MAKE/Makefile.linux
```

If you are not 64-bit (or Apple)

```
gpu_SYSPATH = -L/usr/local/cuda/lib
```

If you are using Apple, compile LAMMPS 32-bit to link with GPU library

```
CC = g++ -m32  
LINK = g++ -m32
```

```
make clean
```

# 6. Add GPU Package to LAMMPS

```
cd $LROOT/src  
make yes-asphere  
make yes-gpu  
make linux
```

# 7. Modify your input script

```
cd $LROOT/bench  
emacs in.lj
```

Must add ‘newton off’ to beginning of script and ‘/gpu’ to a supported pair\_style

```
newton off  
  
...  
  
pair_style lj/cut/gpu one/node 0 2.5
```

GPU Selection Keyword      GPU ID

# 7. Modify your input script (2)

- GPU Selection Keyword
  - **one/node** - single compute "node", which may have multiple cores and/or GPUs. *GpuID* should be set to the ID of the (first) GPU you wish to use with LAMMPS
  - **one/gpu** - multiple compute "nodes" with one GPU per node. *GpuID* should be set to the ID of the GPU.
  - **multi/gpu** - multiple compute "nodes" on your system with multiple GPUs. *GpuID* should be set to the number of GPUs per node

# 8. Run your input script

- **Number of procs = number of gpus you want**

```
mpirun -np 3 lmp_linux < in.lj
```

```
-----  
- Using GPGPU acceleration for LJ-Cut:  
-----  
GPU 1: Tesla C1060, 240 cores, 4 GB, 1.3 GHZ  
GPU 2: Tesla C1060, 240 cores, 4 GB, 1.3 GHZ  
GPU 3: GeForce GTX 295, 240 cores, 0.87 GB, 1.2 GHZ  
-----
```

```
-----  
GPU Time Stamps:  
-----  
Atom copy:      0.07111 s.  
Neighbor copy: 0.0004615 s.  
LJ calc:        0.1702 s.  
Answer copy:    0 s.  
-----
```

# 9. Speed-ups

- Depends on
  - Your CPU
  - Your GPU
  - Number of Particles
  - Cutoff
- More talks showing the GPU acceleration in LAMMPS to come...

# Questions